



Neuroscience of Learning, Memory and Cognition

Instructor: Prof. Karbalaei

Sharif University of Technology

Project: Solving Maze Using Q-Learning Algorithm

By Reza Nayeb Habib
Student ID# 401102694

Contents

1	The Maze Generation Algorithm	1
2	Drawing The Maze	2
3	Q-Learning Algorithm	2
3.1	Full Code and Results	2
3.2	Code Explanation	8
3.3	Deep Q-Learning Implementation	9

1 The Maze Generation Algorithm

The Maze Generation Algorithm is coded as below

```

1 # Adapted from http://code.activestate.com/recipes/578356-random-maze-generator/
2 # Random Maze Generator using Depth-first Search
3 # http://en.wikipedia.org/wiki/Maze_generation_algorithm
4
5
6 import random
7 import matplotlib.pyplot as plt
8 import numpy as np
9
10 random.seed(401102694)
11 mx = 20; my = 20 # width and height of the maze
12
13 maze = [[0 for x in range(mx)] for y in range(my)]
14 dx = [0, 1, 0, -1]; dy = [-1, 0, 1, 0] # 4 directions to move in the maze
15 color = [(0, 0, 0), (255, 255, 255)] # RGB colors of the maze
16
17 # start the maze from a random cell
18 cx = random.randint(0, mx - 1)
19 cy = random.randint(0, my - 1)
20 maze[cy][cx] = 1
21 stack = [(cx, cy, 0)] # stack element: (x, y, direction)
22
23 while len(stack) > 0:
24     (cx, cy, cd) = stack[-1]
25     # to prevent zigzags:
26     # if changed direction in the last move then cannot change again
27     if len(stack) > 2:
28         if cd != stack[-2][2]: dirRange = [cd]
29         else: dirRange = range(4)
30     else: dirRange = range(4)
31
32     # find a new cell to add
33     nlst = [] # list of available neighbors
34     for i in dirRange:
35         nx = cx + dx[i]
36         ny = cy + dy[i]
37         if nx >= 0 and nx < mx and ny >= 0 and ny < my:
38             if maze[ny][nx] == 0:
39                 ctr = 0 # of occupied neighbors must be 1
40                 for j in range(4):
41                     ex = nx + dx[j]; ey = ny + dy[j]
42                     if ex >= 0 and ex < mx and ey >= 0 and ey < my:
43                         if maze[ey][ex] == 1: ctr += 1
44                 if ctr == 1: nlst.append(i)
45
46     # if 1 or more neighbors available then randomly select one and move
47     if len(nlst) > 0:
48         ir = nlst[random.randint(0, len(nlst) - 1)]
49         cx += dx[ir]; cy += dy[ir]; maze[cy][cx] = 1
50         stack.append((cx, cy, ir))
51     else: stack.pop()
52
53 maze = np.array(maze)
54 maze -= 1
55 maze = abs(maze)
56
57 maze[0][0] = 0
58 maze[mx-1][my-1] = 0

```

```
59
60 np.save('maze', np.array(maze))
```

In the project document it is stated that there is always a path between its upper left point (0,0) to its lower right point (my-1,mx-1) this statement is true because the maze Generation Algorithm starts at a random point in the plane then carves out into different paths by seeing the blocks that have more than 2 wall neighbors (0 in the code) and this algorithm always uses already created path points neighbors to create new path points thus the maze is always connected. At the final lines, this algorithm adds (0,0) and (my-1,mx-1) to the path; but how does it make sure that these two points will be also connected by the previous path?

We can ensure these two points are always connected because no cell in this structure will have more than 2 walls beside it so it means every wall will always have at least one path neighbor! Thus when for example (0,0) cell is set to 1 it will always have a neighbor with value one and it will not be an isolated path point and it will be connected to all other path points.

2 Drawing The Maze

The following code draws the Maze by Loading the file and using matplotlib's Imshow:

```
1 maze = np.load('maze.npy')
2 plt.imshow(1-maze, cmap='gray') # black is wall white is path
3 plt.xticks([]); plt.yticks([])
4 plt.show()
```

it gives the following result:




Figure 1: the unique maze of my student number seed

3 Q-Learning Algorithm

3.1 Full Code and Results

```
1 import numpy as np
2 import random
3 import matplotlib.pyplot as plt
4 import imageio
5 import os
```

```

6
7 class Environment:
8     def __init__(self, maze_file):
9         self.maze = np.load(maze_file)
10        self.ny, self.nx = self.maze.shape
11        self.start = (0, 0)
12        self.goal = (self.nx - 1, self.ny - 1)
13        self.maze_file = maze_file
14        self.output_dir = os.path.join(os.path.dirname(maze_file), "output")
15        os.makedirs(self.output_dir, exist_ok=True)
16        # Create output directory if it doesn't exist
17
18    def get_valid_actions(self, x, y):
19        """Returns a list of valid actions for a given state."""
20        valid_actions = []
21        actions = [(0, -1), (1, 0), (0, 1), (-1, 0)] # Up, Right, Down, Left
22        for i, (dx, dy) in enumerate(actions):
23            nx_, ny_ = x + dx, y + dy
24            if 0 <= nx_ < self.nx and 0 <= ny_ < self.ny and self.maze[ny_][nx_] == 0:
25                valid_actions.append(i)
26        return valid_actions
27
28    def get_reward(self, x, y):
29        """Returns the reward for a given state."""
30        if (x, y) == self.goal:
31            return 100 # Big reward for reaching the goal
32        else:
33            return -1 # Small penalty for each step
34
35    def is_goal(self, x, y):
36        """Checks if the given state is the goal."""
37        return (x, y) == self.goal
38
39    def visualize_maze(self):
40        """Visualizes the maze."""
41        plt.imshow(self.maze, cmap='gray')
42        plt.title("Maze")
43        plt.show()
44
45 class Agent:
46     def __init__(self, env, alpha=0.1, gamma=0.9, epsilon=1.0,
47                  epsilon_decay=0.995, epsilon_min=0.01):
48         self.env = env
49         self.alpha = alpha # Learning rate
50         self.gamma = gamma # Discount factor
51         self.epsilon = epsilon # Initial exploration rate
52         self.epsilon_decay = epsilon_decay # Decay per episode
53         self.epsilon_min = epsilon_min # Minimum epsilon value
54         self.actions = [(0, -1), (1, 0), (0, 1), (-1, 0)] # Up, Right, Down, Left
55         self.num_actions = len(self.actions)
56         self.Q_table = np.zeros((env.ny, env.nx, self.num_actions)) # Q-table
57
58    def choose_action(self, x, y):
59        """Epsilon-greedy policy."""
60        if random.uniform(0, 1) < self.epsilon:
61            return random.choice(self.env.get_valid_actions(x, y)) # Explore
62        else:
63            valid_actions = self.env.get_valid_actions(x, y)
64            return max(valid_actions, key=lambda a: self.Q_table[y, x, a]) # Exploit
65
66    def train(self, num_episodes):
67        """Trains the agent using Q-learning."""

```

```

68     for episode in range(num_episodes):
69         x, y = self.env.start # Start position
70         total_reward = 0
71
72         while not self.env.is_goal(x, y): # Until reaching the goal
73             action = self.choose_action(x, y)
74             dx, dy = self.actions[action]
75             new_x, new_y = x + dx, y + dy
76
77             # Reward system
78             reward = self.env.get_reward(new_x, new_y)
79
80             # Q-value update
81             best_next_action = max(self.env.get_valid_actions(new_x, new_y),
82                                   key=lambda a: self.Q_table[new_y, new_x, a])
83             self.Q_table[y, x, action] += self.alpha * (reward + self.gamma *
84                                                         self.Q_table[new_y, new_x, best_next_action] - self.Q_table[y, x, action])
85
86             x, y = new_x, new_y
87             total_reward += reward
88
89             # Decay epsilon
90             self.epsilon = max(self.epsilon * self.epsilon_decay, self.epsilon_min)
91
92             if episode % 500 == 0:
93                 print(f"Episode {episode}, Total Reward:
94                       {total_reward}, Epsilon: {self.epsilon:.3f}")
95
96     def solve_maze(self):
97         """Uses the trained Q-table to solve the maze."""
98         x, y = self.env.start
99         path = [(x, y)]
100         while not self.env.is_goal(x, y):
101             action = max(self.env.get_valid_actions(x, y),
102                           key=lambda a: self.Q_table[y, x, a])
103             dx, dy = self.actions[action]
104             x, y = x + dx, y + dy
105             path.append((x, y))
106         return path
107
108     def draw_policy(self, episode, save_path=None):
109         """Visualizes the best policy at each cell using arrows.
110         The optimal path (even if it leads to dead ends) is shown in green.
111         """
112         fig, ax = plt.subplots(figsize=(10, 10))
113         ax.imshow(1 - self.env.maze, cmap="gray") # Show maze
114
115         arrow_map = {
116             0: (0, -1), # Up
117             1: (1, 0), # Right
118             2: (0, 1), # Down
119             3: (-1, 0) # Left
120         }
121
122         # Compute the best current path, even if not optimal (can hit dead ends)
123         x, y = self.env.start
124         visited = set()
125         best_path = []
126
127         while (x, y) not in visited and not self.env.is_goal(x, y):
128             visited.add((x, y))
129             best_path.append((x, y))

```

```

130         valid_actions = self.env.get_valid_actions(x, y)
131         if not valid_actions: # If no valid moves, stop
132             break
133
134
135         best_action = max(valid_actions, key=lambda a: self.Q_table[y, x, a])
136         dx, dy = arrow_map[best_action]
137         x, y = x + dx, y + dy
138
139     best_path_set = set(best_path) # Convert to set for quick lookup
140
141     # Draw policy arrows
142     for y in range(self.env.ny):
143         for x in range(self.env.nx):
144             if self.env.maze[y, x] == 0: # Ignore walls
145                 valid_actions = self.env.get_valid_actions(x, y)
146                 if valid_actions: # If there are valid actions
147                     best_action = max(valid_actions, key=lambda a: self.Q_table[y, x, a])
148                     dx, dy = arrow_map[best_action]
149
150                     color = 'green' if (x, y) in best_path_set else 'red'
151                     ax.arrow(x, y, dx * 0.3, dy * 0.3, head_width=0.2,
152                             head_length=0.2, fc=color, ec=color)
153
154     ax.set_title(f"Policy Visualization - Episode {episode}")
155
156     if save_path:
157         plt.savefig(save_path)
158         plt.close(fig) # Close the figure to avoid displaying it
159     else:
160         plt.show()
161
162
163 def draw_optimal_path(self, path=None, save_path=None):
164     """Draws the maze with the optimal path overlaid and optionally saves it as a PNG."""
165     if path is None:
166         path = self.solve_maze() # Solve the maze if no path is provided
167
168     fig, ax = plt.subplots(figsize=(10, 10))
169     ax.imshow(self.env.maze, cmap="gray") # Show maze
170
171     # Extract X and Y coordinates from path
172     x_coords, y_coords = zip(*path)
173
174     # Plot the path with red line
175     ax.plot(x_coords, y_coords, color='red', linewidth=2, marker='o'
176            , markersize=4, label="Optimal Path")
177
178     # Mark start and goal positions
179     ax.scatter([self.env.start[0]], [self.env.start[1]], color='green'
180            , s=100, label="Start (0,0)")
181     ax.scatter([self.env.goal[0]], [self.env.goal[1]], color='blue'
182            , s=100, label=f"Goal ({self.env.goal[0]},{self.env.goal[1]})")
183
184     ax.legend()
185     plt.title("Optimal Path in the Maze")
186
187     if save_path:
188         plt.savefig(save_path)
189         plt.close(fig) # Close the figure to avoid displaying it
190     else:
191         plt.show()

```

```

192
193 def train_with_visualization(self, num_episodes, frame_time, d_episode):
194     """Trains the agent and saves a GIF of the policy evolution."""
195     frames = [] # List to store frames
196
197     for episode in range(num_episodes):
198         x, y = self.env.start # Start position
199         total_reward = 0
200
201         while not self.env.is_goal(x, y): # Until reaching the goal
202             action = self.choose_action(x, y)
203             dx, dy = self.actions[action]
204             new_x, new_y = x + dx, y + dy
205
206             # Reward system
207             reward = self.env.get_reward(new_x, new_y)
208
209             # Q-value update
210             best_next_action = max(self.env.get_valid_actions(new_x, new_y),
211                                   key=lambda a: self.Q_table[new_y, new_x, a])
212             self.Q_table[y, x, action] += self.alpha * (reward + self.gamma *
213                                                         self.Q_table[new_y, new_x, best_next_action] - self.Q_table[y, x, action])
214
215             x, y = new_x, new_y
216             total_reward += reward
217
218             # Decay epsilon
219             self.epsilon = max(self.epsilon * self.epsilon_decay, self.epsilon_min)
220
221             # Save policy visualization at key episodes
222             if episode % d_episode == 0 or episode == num_episodes - 1:
223                 save_path = os.path.join(self.env.output_dir, f"policy_episode_{episode}.png")
224                 self.draw_policy(episode, save_path)
225                 frames.append(imageio.imread(save_path)) # Store frame
226
227             if episode % d_episode == 0:
228                 print(f"Episode {episode}, Total Reward: {total_reward},
229                       Epsilon: {self.epsilon:.3f}")
230
231             # Save GIF with longer duration between frames
232             gif_path = os.path.join(self.env.output_dir, "policy_evolution.gif")
233             imageio.mimsave(gif_path, frames, duration=frame_time*num_episodes/d_episode)
234             print(f"GIF saved as {gif_path}")
235
236             # Clean up temporary files
237             for episode in range(0, num_episodes, d_episode):
238                 os.remove(os.path.join(self.env.output_dir, f"policy_episode_{episode}.png"))
239                 os.remove(os.path.join(self.env.output_dir, f"policy_episode_{num_episodes - 1}.png"))
240
241
242 # Main execution
243 if __name__ == "__main__":
244     env = Environment('.\\maze.npy')
245     agent = Agent(env)
246
247     # Train the agent
248     agent.train_with_visualization(num_episodes=1000, frame_time=5.0, d_episode=25)
249
250     # Solve the maze and visualize the optimal path
251     solution_path = agent.solve_maze()
252
253     optimal_path_image_path = os.path.join(env.output_dir, "optimal_path.png")

```



```

254 agent.draw_optimal_path(solution_path, save_path=optimal_path_image_path)
255 print(f"Optimal path visualization saved as {optimal_path_image_path}")
256 agent.draw_optimal_path(solution_path)
257

```

This code gives the following Results (the Gif file is included in the zip sent with this pdf file):




Figure 2: the Optimal path in the maze found by Q-learning




Figure 3: the Optimal policy in the maze found by Q-learning

3.2 Code Explanation

Environment:

This class contains the maze cells mapping the reward function of the maze and the path the maze file is loaded from. Its variables contain the goal point of the maze, the start point of the maze, the maze shape and ...

`get_valid_action(x,y):`

This method of the `Environment` class gives the valid actions that can be taken at position (x,y) without hitting a wall or going outside of the maze.

`get_reward(x,y):`

This method gives the reward of any given cell (x,y). all cells have a default reward of -1 to avoid the agent from wandering and punish him for taking longer paths. Also the goal point has a reward of 100 for motivating the agent to reach the end point.

`is_goal(x,y):`

Checks if the cell is the goal cell; useful in functions needing to check if they have reached the end of the maze or not. Returns True if (nx,ny) is given.

`visualize_maze()`:

Uses `Imshow` to draw the maze.

Agent:

This class creates an agent that uses Q-Learning (Explained Later) to find the best policy for traversing in the maze. This class has the Environment, learning rate, Discount Factor, exploration rate and its decay factor, and the Q-table as its variables. It creates Q-table from the size of the maze and the number of actions that can be taken (number of directions that can be traversed) and Initializes them as zeros.

`choose_action(x,y)`:

This method of the Agent class is essentially the actor in actor-critic algorithm. It uses the Q-table to choose the best policy and act on that policy, this action is called exploitation; It also does exploration by generating a random value between 0 and 1 and if the value is smaller than the exploration probability (epsilon) it chooses its action randomly from valid actions.

`train(num_episodes)` and `train_with_visualization(num_episodes, frame_time, d_episode)`:

These functions implement the formula for Q-learning as stated below:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \quad (1)$$

by this line of code:

```
1 self.Q_table[y, x, action] += self.alpha * (reward + self.gamma
2   * self.Q_table[new_y, new_x, best_next_action] - self.Q_table[y, x, action])
```

In every iteration it chooses an action by `choose_action(x,y)` and then updates the Q-table according the action taken. This process is repeated until the goal state is reached, then the episode is done and a new one starts it iterates this until the number of given episodes of learning is reached. The function with visualize is the same it only Visualizes the best policy using `draw_policy()` and saves it periodically on specific period of episodes; after that it creates a gif from the given files and saves it in the output folder.

`draw_policy(episode, save_path=None)`

This function first draws the maze with `Imshow` then uses `scatter` to draw the best policy at each cell with an arrow (it draws the arrows on the optimal path with green and others with red).

`draw_optimal_path (path=None, save_path =None):`

This function finds the optimal path by using `solve_maze()` and getting the path that will be taken on exploitation mode and again uses `scatter` for drawing the optimal path given by `solve_maze`.

`solve_maze()`:

This fuction follows the path given by Q-table until it reaches the goal point. It should not be used before the model is completely trained to reach the goal point otherwise it will enter an infinite loop.

3.3 Deep Q-Learning Implementation

The Deep Q-Learning algorithm is very similar to normal Q-Learning. The difference is that the state variables (here x and y) are fed into a deep neural network and it gives the Q value for all possible actions. Here we show how we implemented this with pytorch:

```
1 import numpy as np
2 import random
```

```

3 import torch
4 import torch.nn as nn
5 import torch.optim as optim
6 import torch.nn.functional as F
7 import matplotlib.pyplot as plt
8 import imageio
9 import os
10 from collections import deque
11
12 # Environment Class
13 class Environment:
14     def __init__(self, maze_file):
15         self.maze = np.load(maze_file)
16         self.ny, self.nx = self.maze.shape
17         self.start = (0, 0)
18         self.goal = (self.nx - 1, self.ny - 1)
19         self.actions = [(0, -1), (1, 0), (0, 1), (-1, 0)] # Up, Right, Down, Left
20         self.output_dir = os.path.join(os.path.dirname(maze_file), "output_dqn")
21         os.makedirs(self.output_dir, exist_ok=True)
22
23     def get_valid_actions(self, x, y):
24         """Returns a list of valid actions for a given state."""
25         valid_actions = []
26         for i, (dx, dy) in enumerate(self.actions):
27             nx_, ny_ = x + dx, y + dy
28             if 0 <= nx_ < self.nx and 0 <= ny_ < self.ny and self.maze[ny_, nx_] == 0:
29                 valid_actions.append(i)
30         return valid_actions
31
32     def get_reward(self, x, y):
33         """Returns the reward for a given state."""
34         return 100 if (x, y) == self.goal else -1
35
36     def is_goal(self, x, y):
37         """Checks if the given state is the goal."""
38         return (x, y) == self.goal
39
40 # Deep Q-Network (DQN) Model
41 class DQN(nn.Module):
42     def __init__(self, input_dim, output_dim):
43         super(DQN, self).__init__()
44         self.fc1 = nn.Linear(input_dim, 128)
45         self.fc2 = nn.Linear(128, 128)
46         self.fc3 = nn.Linear(128, output_dim)
47
48     def forward(self, x):
49         x = F.relu(self.fc1(x))
50         x = F.relu(self.fc2(x))
51         return self.fc3(x)
52
53 # DQN Agent
54 class DQNAgent:
55     def __init__(self, env, gamma=0.9, lr=0.001, batch_size=64, memory_size=50000):
56         self.env = env
57         self.gamma = gamma
58         self.batch_size = batch_size
59         self.epsilon = 1.0
60         self.epsilon_decay = 0.995
61         self.epsilon_min = 0.01
62         self.memory = deque(maxlen=memory_size)
63
64         self.model = DQN(2, len(env.actions))

```

```

65     self.target_model = DQN(2, len(env.actions))
66     self.target_model.load_state_dict(self.model.state_dict())
67     self.optimizer = optim.Adam(self.model.parameters(), lr=lr)
68
69     def get_state(self, x, y):
70         """Encodes state as a tensor."""
71         return torch.tensor([x / self.env.nx, y / self.env.ny], dtype=torch.float32)
72
73     def choose_action(self, x, y):
74         """Epsilon-greedy policy."""
75         if random.random() < self.epsilon:
76             return random.choice(self.env.get_valid_actions(x, y))
77         else:
78             state = self.get_state(x, y).unsqueeze(0)
79             q_values = self.model(state)
80             valid_actions = self.env.get_valid_actions(x, y)
81             return max(valid_actions, key=lambda a: q_values[0, a].item())
82
83     def store_experience(self, state, action, reward, next_state, done):
84         """Stores experience in replay memory."""
85         self.memory.append((state, action, reward, next_state, done))
86
87     def train(self):
88         """Trains the model using replay memory."""
89         if len(self.memory) < self.batch_size:
90             return
91
92         batch = random.sample(self.memory, self.batch_size)
93         states, actions, rewards, next_states, dones = zip(*batch)
94
95         states = torch.stack(states)
96         next_states = torch.stack(next_states)
97         actions = torch.tensor(actions, dtype=torch.long)
98         rewards = torch.tensor(rewards, dtype=torch.float32)
99         dones = torch.tensor(dones, dtype=torch.float32)
100
101         q_values = self.model(states).gather(1, actions.unsqueeze(1)).squeeze(1)
102         next_q_values = self.target_model(next_states).max(1)[0].detach()
103         target_q_values = rewards + self.gamma * next_q_values * (1 - dones)
104
105         loss = F.mse_loss(q_values, target_q_values)
106         self.optimizer.zero_grad()
107         loss.backward()
108         self.optimizer.step()
109
110     def update_target_network(self):
111         """Updates target network weights."""
112         self.target_model.load_state_dict(self.model.state_dict())
113
114     def train_agent(self, num_episodes, update_target_every=50, frame_time=5.0, d_episode=50):
115         """Trains the DQN agent and creates a GIF."""
116         frames = []
117         for episode in range(num_episodes):
118             x, y = self.env.start
119             total_reward = 0
120             state = self.get_state(x, y)
121
122             while not self.env.is_goal(x, y):
123                 action = self.choose_action(x, y)
124                 dx, dy = self.env.actions[action]
125                 new_x, new_y = x + dx, y + dy

```

```

127         reward = self.env.get_reward(new_x, new_y)
128         next_state = self.get_state(new_x, new_y)
129         done = self.env.is_goal(new_x, new_y)
130
131         self.store_experience(state, action, reward, next_state, done)
132         self.train()
133
134         x, y = new_x, new_y
135         state = next_state
136         total_reward += reward
137
138         self.epsilon = max(self.epsilon * self.epsilon_decay, self.epsilon_min)
139
140         if episode % update_target_every == 0:
141             self.update_target_network()
142
143         if episode % d_episode == 0:
144             save_path = os.path.join(self.env.output_dir, f"policy_{episode}.png")
145             self.visualize_policy(save_path)
146             frames.append(imageio.imread(save_path))
147
148             print(f"Episode {episode}, Total Reward: {total_reward}, Epsilon: {self.epsilon:.3f}")
149
150         gif_path = os.path.join(self.env.output_dir, "policy_evolution.gif")
151         imageio.mimsave(gif_path, frames, duration=frame_time*num_episodes/d_episode)
152         print(f"GIF saved as {gif_path}")
153
154     def visualize_policy(self, save_path=None):
155         """Visualizes the current policy."""
156         fig, ax = plt.subplots(figsize=(10, 10))
157         ax.imshow(1 - self.env.maze, cmap="gray")
158
159         for y in range(self.env.ny):
160             for x in range(self.env.nx):
161                 if self.env.maze[y, x] == 0:
162                     best_action = max(self.env.get_valid_actions(x, y), key=lambda a: self.model(self.ge
163                     dx, dy = self.env.actions[best_action]
164                     ax.arrow(x, y, dx * 0.3, dy * 0.3, head_width=0.2, head_length=0.2, fc="red", ec="re
165
166         if save_path:
167             plt.savefig(save_path)
168             plt.close()
169         else:
170             plt.show()
171
172 # Main Execution
173 if __name__ == "__main__":
174     env = Environment("maze100.npy")
175     agent = DQNAgent(env)
176     agent.train_agent(num_episodes=5000)
177
178 # Save final policy visualization
179 final_policy_path = os.path.join(env.output_dir, "final_policy.png")
180 agent.visualize_policy(final_policy_path)

```